

Mojo-V ISA Extension Specification

Document version: 0.90, Mojo-V specification version (`mojov_ver`) #0x1

Release date: October 24, 2025.

To provide feedback: Email: mojov-devs@umich.edu

NOTE: The latest version of this document is here: <https://github.com/toddmaustin/mojo-v>

1. Introduction

Mojo-V (pronounced “mojo-five”) extends the RISC-V RV64 ISA to support **secret computation** — computation that is **blind** to developers and **silent** in that its behavior does not reveal secret values. It provides cryptographic-strength privacy with minimal performance overhead and enables processors to execute operations on encrypted data without exposing values to any software, developer, or IT operator.

Mojo-V enforces privacy entirely in hardware, rendering all software and developers untrusted. Every software-visible operation is either data-oblivious or cryptographically sealed within the processor.

Developed at the University of Michigan and Agita Labs, Mojo-V is the fifth generation in a lineage of secure architectures combining strong cryptography with secure hardware to advance system security and privacy:

Morpheus (I) → Morpheus II (II) → Sequestered Encryption Unit (III) → TrustForge Platform (IV) → Mojo-V (V)

2. Caveats

Mojo-V is an open project that is in progress. As such, not everything is finished, either because we don’t have the resources yet or we are waiting for community input. So, please be patient and please contribute. Here are current caveats with this specification:

- Currently, Mojo-V is only specified for extending RV64 RISC-V. Would an RV32 extension be desired?
- CSR register addresses and lde/sde opcode assignments are not finalized as yet, as we will need to work with the RISC-V community to finalize these values.
- The hardware ciphers are not yet specified. Currently, we only specify the mechanism used to probe the hardware to see what ciphers it supports and the data contract field where the data owner can specify the ciphers they require to protect their data. We expect to codify this part of the Mojo-V specification once real hardware implementations

start to materialize. As of this writing, the open-source project uses ML-KEM (Kyber) and Simon-128 for its asymmetric and symmetric key ciphers, respectively.

3. Secret Computation

As illustrated in **Figure 1**, Mojo-V enables privacy-preserving computation by ensuring that even untrusted cloud or edge infrastructure can process encrypted data safely and efficiently. Imagine a **smart-home security camera** that performs face and object recognition using an AI model hosted in the cloud.

Without Mojo-V, the cloud service must first decrypt each video frame to analyze it. The decrypted video is accessible to the service provider's software, developers, or IT staff, meaning privacy depends entirely on the provider's internal trustworthiness.

With Mojo-V, the camera encrypts each video frame using the data owner's key and transmits both the encrypted data and a *third-party encrypted data contract* to a Mojo-V-enabled cloud server. The Mojo-V processor decrypts and analyzes the data entirely **within hardware**, so plaintext video is never visible to software or operators. Even a compromised operating system or malicious programmer cannot view, log, or leak the contents. With Mojo-V, **trust shifts completely into hardware**, removing the need to rely on software or developers for data privacy.




Figure 1. Privacy-Enhanced Cloud AI Application using Mojo-V.

Mojo-V implements **blind computation** through a hardware mechanism that isolates and controls the flow of secret data. It introduces the concept of **secret registers** and specialized semantics that strictly confine secrets within the processor. Data may only **enter** secret registers from a **third-party encrypted load**, and once marked secret, **every derived result** automatically inherits secrecy, forcing it back into a secret register. The **only legal path** for data to **leave** the secret registers via a **third-party encrypted store**, which re-encrypts the value under the data owner's key using strong, authenticated encryption. This design effectively **sequesters** secret values within hardware. Whenever data is written to memory, it is encrypted under a key controlled by the data owner, never by the service provider. Authenticated encryption ensures that any tampering or manipulation attempts are immediately detected by Mojo-V hardware during decryption.

Mojo-V also enforces **silent computation**, preventing any program behavior from revealing secret information. Hardware rules guarantee that program execution is completely data-oblivious — no secret value can influence control flow, operational timing, or memory access patterns. Three hardware-enforced restrictions make this possible:

1. **No secret registers in branch predicates** — prevents attackers from deducing secrets through timing differences, speculative execution, or instruction-path divergence.
2. **No secret code pointers or return addresses** — ensures that secret data cannot direct control flow or alter instruction fetch sequences, blocking instruction-cache and branch predictor side channels.
3. **No secret memory addresses** — prevents secret-dependent data memory references that could leak information through cache behavior, TLB lookups, or page faults.

These constraints eliminate observable differences between executions that depend on secret data. Every instruction sequence behaves identically from an external perspective, regardless of the secrets being processed.

Even under these restrictions, any Mojo-V computation must be expressed with a **data-oblivious programming model**. Traditional control flow and indexing can be restructured using conditional moves, oblivious lookup tables, or fixed-access patterns. The programmer writes normal code, while Mojo-V's ISA and compiler toolchain ensure that operations on secret data respect the enforced rules. Privacy is achieved automatically through hardware enforcement. Any program can be made to be data oblivious, and only code that computes on sensitive data needs to be this way.

By combining **blind** and **silent** computation, Mojo-V establishes a framework for truly private, hardware-enforced computation. Secret data remains confined to secret registers and encrypted memory, while all observable software behavior remains independent of the secrets being processed. With Mojo-V, **privacy is guaranteed by silicon itself**, not by software policy or developer discipline.

4. Mojo-V Security Model

A **security model** defines what must be protected, who or what can be trusted, and how an attacker is expected to behave. It specifies the assumptions about the system, the capabilities and limits of adversaries, and what constitutes a successful breach. For Mojo-V, the model is intentionally conservative — granting attackers wide control over software and system resources while relying solely on hardware enforcement for secrecy and integrity.

Mojo-V follows a **zero-trust architecture**, assuming that all software — including operating systems, hypervisors, programmers, and administrators — may be compromised. Only Mojo-V hardware and the cryptographically authenticated data-owner channels are considered trustworthy. Table 1 depicts these trust boundaries, showing how Mojo-V isolates secret data inside hardware-protected domains while exposing only ciphertext to untrusted software layers.

In this model, only Mojo-V hardware may compute on decrypted data. All other entities interact exclusively with ciphertext. Any violation of this trust boundary triggers a Mojo-V security exception, terminating execution before a secret can be leaked.

Table 1. Mojo-V Security Model — Trust Boundaries.

Entity	Trust Level	Rationale
Software	Untrusted	Software compromise cannot reveal or alter protected data.
Programmers	Untrusted	Malicious application logic cannot intentionally or accidentally expose secrets.
IT / OS Staff	Untrusted	Administrative privileges provide no access to secret registers or encrypted memory.
Hardware (non-secret)	Untrusted	Non-Mojo-V hardware only handles ciphertext or public data.
Hardware (Mojo-V)	Trusted	Performs all secret computation on decrypted data inside sealed hardware registers.

Mojo-V's **assumptions** include a trusted hardware root of trust that is free of backdoors and design errors. Cryptographically authenticated communication channels ensure that only the data owner can install or refresh data contracts. A hardware-based true random number generator (TRNG) supplies salts for encryption freshness, while the physical silicon is assumed to be fabricated without malicious tampering.

The **attacker** in Mojo-V's model is extraordinarily powerful. The adversary can control every layer of system software, execute arbitrary code at any privilege level, and even possess root access to servers. They may attempt side-channel analysis via operational timing, cache contention, branch prediction priming, or port contention analysis. Yet their influence ends at the hardware boundary. Attackers cannot read secret registers, extract plaintext through debug ports (e.g., JTAG), or forge ciphertext without valid keys and authentication hashes. They also cannot manipulate ciphertext undetected, as authenticated encryption ensures integrity and replay protection.

An **attack succeeds** only if an attacker achieves one of three outcomes:

1. **Secret disclosure** – obtaining plaintext or key material.
2. **Ciphertext Forgery** – introducing falsified ciphertext that passes authentication.
3. **Replay compromise** – reusing old ciphertext to alter a computation.

Mojo-V's design prevents all three through hardware-enforced confidentiality, integrity, and freshness. Even with total software compromise, the system guarantees that secrets cannot be revealed, forged, or replayed. The result is a particularly powerful and resilient threat model — one that concedes near-total control to the attacker, yet still upholds the fundamental promise of Mojo-V: **no unwanted disclosures, even when every line of software is compromised**.

5. Mojo-V ISA Specification

This section defines the Mojo-V ISA extensions for secret computation on RISC-V.

5.1 Mojo-V Configuration CSRs

The Mojo-V Configuration Control and Status Registers (CSRs) define the architectural interface between hardware secrecy mechanisms and system-level software. These registers establish and maintain the cryptographic contract between the processor and the data owner, manage key installation, and track configuration state throughout secret computation. Together, they provide the necessary state control to enforce Mojo-V's zero-trust execution model, where secrecy is guaranteed entirely by hardware.

Table 2. Secret Configuration Registers.

CSR Name	Address	Access	Purpose
<code>mojov_cfg</code>	0x0A0	R/W	Controls secret configuration enable (<code>mojov_en</code>), key validity, encryption format, and zeroization state.
<code>mojov_ciphers</code>	0x0A1	Read-only	64-bit value indicating the asymmetric and symmetric key ciphers that are supported by this Mojo-V implementation. [Format TBD]

<code>mojov_pubkey</code>	0x0A2	Read-only	Provides the processor's public key for data owner attestation. [Format TBD]
<code>mojov_keycfg</code>	0x0A3	Write-only	Installs a 512-bit encrypted data contract defining the data owner's key, permitted operations, and embedded authentication hash (<code>auth_hash</code>) used for freshness and replay protection.
<code>mojov_keystate</code>	0x0A4	R/W	Stores encrypted key state for context save/restore operations. [Format TBD]

The `mojov_cfg` register is the central control register for Mojo-V secret computation. It governs whether the processor is operating in secret mode, indicates if a valid data contract is installed, and records the active encryption format and implementation version. It can be accessed and altered in User mode by any OS or application program. The `mojov_cfg` register contains the following bitfields:

Table 3. Bitfield Layout of the Mojo-V CSR Configuration Register. The `mojov_cfg` register controls activation of secret computation and tracks key validity, encryption format, and implementation version for Mojo-V's operation.

```
| 31          11|10   3|2|1|0|
| Reserved | VER | F|K|E|
```

Bit	Field	Abbrev	Access	Description
0	<code>mojov_en</code>	E	R/W	Enables secret computation mode.
1	<code>key_valid</code>	K	Read-only	Indicates a valid data contract is installed.
2	<code>format_sel</code>	F	Read-only	Active encrypted memory format (0 = weak, 1 = strong). Read-only.
3	<code>mojov_ver</code>	VER	Read-only	Contains an 8-bit version ID for the Mojo-V implementation, ensuring backward compatibility across revisions.

When `mojov_en` is asserted, Mojo-V enforces the secret computation rules described in Section 2, limiting data movement between secret and public domains. The `key_valid` bit signals that a legitimate data contract is active; operations involving secret data cannot proceed unless this bit is set. The `format_sel` bit reflects which encryption format (weak or strong) is active, determined by the data-owner's data contract. Finally, the `mojov_ver` field specifies the

hardware implementation version of the Mojo-V ISA extension, allowing compilers and runtime systems to adapt to future revisions of the specification.

The **mojov_ciphers** register is a read-only register that enumerates the asymmetric and symmetric ciphers supported by the current Mojo-V implementation. Each bit corresponds to a specific cipher algorithm; bits that are set to “1” indicate that the cipher is implemented and available for use by hardware cryptographic operations.

The register encodes two cipher groups:

- **Asymmetric key ciphers** — These bits identify public-key algorithms supported by the processor, such as **RSA-2048** and **Elliptic Curve P-256**.
- **Symmetric key ciphers** — These bits identify data encryption algorithms available for third-party encryption, such as **AES-128/256** and **Simon-128**.

Software or remote attestation tools may read **mojov_ciphers** to determine which cryptographic primitives are supported by a particular Mojo-V device before establishing a data contract or provisioning encryption keys. Each bit position is implementation-defined, and future revisions of Mojo-V may extend this register to indicate additional cipher suites or key-derivation mechanisms.

The **mojov_pubkey** register exposes a processor-resident public key used to verify authenticity to data owners and establish data contracts. This key enables cryptographic attestation that the processor is a genuine Mojo-V implementation running approved microcode. The detailed structure and provisioning method for **mojov_pubkey** will be defined in a future version of this specification.

The **mojov_keycfg** register is a write-only register that receives a 256-bit encrypted data contract from the data owner. This contract defines the symmetric or asymmetric keys to be used, permissible operations, encryption formats, and an embedded **auth_hash** to ensure freshness and replay protection. The exact format of the data contract and the associated cryptographic protocol for loading and verifying it will be detailed in a later version of this document.

The **mojov_keystate** register preserves encrypted state across privilege transitions and context switches. When an operating system performs a context save or restore, the **mojov_keystate** register allows secure transfer of key state without exposing plaintext material to software. This ensures that secret computation can resume seamlessly after preemption, with confidentiality and integrity fully preserved.

5.2 Mojo-V Exceptions

Mojo-V introduces a new architectural exception type known as the **Mojo-V Security Exception**, assigned to **exception code 0x1F** (reserved for architecture-specific security

events). This exception is raised whenever a program attempts an operation that could disclose secret data or excite a side channel. All detection and enforcement occur strictly at **decode time**, before any instruction is executed.

A Mojo-V Security Exception is triggered under the following conditions:

- Any instruction attempts to move or expose a secret value to a public register or unencrypted memory location.
- A secret value is used as a branch predicate, load/store address, or code pointer.
- A conversion or move operation (**fmv**, **fcvt**, or equivalent) attempts a **secret**→**non-secret** transfer.
- Any instruction violates the secrecy rules defined in Section 2 of this specification.

This exception type allows Mojo-V hardware to prevent data leakage before it occurs, halting instruction issue for unsafe data flows. Because this enforcement happens during instruction decode, no secret-dependent execution ever takes place.

Additionally, secret-mode instructions **never raise runtime exceptions**. Operations such as **div** or floating-point arithmetic that encounter exceptional conditions (e.g., divide-by-zero, underflow) must **drop the exception** and instead attach the condition as metadata to the resulting secret value. This ensures that secret data cannot influence observable program control flow or timing, preserving data-oblivious semantics throughout secret computation.

5.3 Secret Register Semantics

The **secret register semantics** ensure that values held in secret registers cannot be disclosed — neither directly through unencrypted stores or register moves, nor indirectly through timing, speculation, or any microarchitectural side channel. Secret registers have a single 64-bit data value that exists only within hardware; exceptions on secret operations are dropped rather than reported to software.

The Mojo-V architecture designates **x28–x31** and **f28–f31** as secret-capable registers. When secret mode is enabled (**mojov_en = 1**), these become the **secret general-purpose registers** (**p0–p3**) and **secret floating-point registers** (**pf0–pf3**). When secret mode is disabled, these registers revert to their standard RISC-V roles and are **automatically immediately zeroed** to prevent leakage of any secret data.

During instruction decode, the hardware verifies that secret and public operands are used safely. If any instruction violates these semantics, a *Mojo-V security exception* is raised immediately at decode. The rules are as follows:

- Any operation may compute on secret register values, but if any input operand is secret, then the destination register must also be secret.

- A secret register can be initialized through either an encrypted or non-encrypted load; however, third-party encrypted loads may only deposit their decrypted results into secret registers.
- Secret registers cannot be used as: (i) branch predicate inputs (e.g., `bne` operands), (ii) data address inputs (e.g., base registers for `ld` or `st`), or (iii) code pointer inputs (e.g., target address for `jalr`).
- The `fmv` and `fcvt` instructions, which transfer values between the integer and floating-point register files, are permitted only when data flow remains within the secret domain or when moving non-secret data into a secret register. Specifically, **secret**→**secret** and **non-secret**→**secret** conversions are allowed, while any **secret**→**non-secret** conversion will immediately trigger a *Mojo-V security exception* at decode time.

The functional units that operate on secret registers must also ensure that timing behavior is independent of secret inputs. The system need not employ fully constant-time execution units; rather, it must guarantee that latency variations do not depend on secret data values. This requirement applies only to instructions marked as secret during decode.

When an instruction operating on a secret register encounters an exceptional condition (for example, divide-by-zero or floating-point underflow), the exception is dropped rather than reported to software. The hardware masks the event, preventing it from influencing observable control flow or timing. Dependent computations may continue using the resulting data value, ensuring that exception behavior cannot serve as a side channel.

When the operating system initiates a process context switch (due to an interrupt, exception, or RISC-V `ECALL`), if secret registers are enabled, their state must be preserved using **encrypted stores** to the process context save area. Upon resumption, the system performs encrypted loads to restore the registers. Since Mojo-V's model assumes no software trust, even the OS cannot inspect or modify the plaintext contents of these registers. Any corruption or tampering would be detected when the data owner verifies the computation results.

Finally, the only possible way for data to **exit** the secret registers is through a **third-party encrypted store**. In this operation, the value is encrypted under the data owner's key using authenticated encryption before leaving the hardware boundary. This mechanism guarantees that secret data is never visible to software or developers and can only be shared with authorized third-party entities through hardware-enforced, cryptographically protected channels.

5.4 Third-Party Encrypted Loads and Stores

Third-party encrypted loads and stores are the **only** mechanisms by which sensitive data from an external data owner may enter or exit Mojo-V's secret domain. These instructions enable secure cooperation between an untrusted service provider and a trusted data owner, allowing encrypted computation to occur while ensuring that plaintext values are never visible to software, developers, or IT personnel.

The two instructions implementing this capability are **lde** (Load Encrypted) and **sde** (Store Encrypted). Both operate on ciphertext that conforms to a data contract provided by the data owner. The encryption employed is strong, semantically secure, and authenticated through an **auth_sig** value embedded in the data contract. This guarantees that tampering, forgery, or replay of ciphertext values can be detected immediately by hardware.

Sensitive third-party data may be written to memory by an external device or process, provided the corresponding cryptographic key has been installed into Mojo-V via the **mojov_keycfg** CSR. Once installed, the LDE instruction can load this data into secret registers for computation, and SDE can store computation results back to memory, maintaining an end-to-end encrypted processing pipeline. At no point in this flow does unencrypted data appear in any software-visible register or memory location.

LDE/SDE Encrypted Data Formats

Mojo-V supports two plaintext block formats for memory-resident ciphertexts, defined by the data owner's data contract. A weak memory format provides less secure and more time/space efficient third-party encrypted stores, while a strong memory format provides significantly more security at the cost of increased storage and computation costs. An attacker with access to a moderately powerful data center would be able to perform cryptanalysis on weak-format encrypted program values, likely being able to recover boolean values within a few days of analysis. Under strong-format encryption, cryptanalysis is likely not possible with today's data centers. Note that the memory format is specified by the data owner only, when their data control is installed (in the **format_sel** field).

Weak Encryption Format

```
struct { // 128-bits in size, 128-bit alignment
    uint64_t val; // register plaintext value
    uint32_t salt; // random salt
    uint32_t auth_sig; // authentication signature (from contract)
} pt_weak;
```

Strong Encryption Format

```
struct { // 256-bits in size, 128-bit alignment
    uint64_t val; // register plaintext value
    uint64_t salt; // random salt
    uint64_t auth_sig; // authentication signature (from contract)
    uint64_t metadata; // target-specific metadata (e.g., device hash, overflow flag)
} pt_strong;
```

The `auth_sig` field is generated under the data owner's data contract to authenticate ciphertext integrity and freshness. The data owner may periodically rotate or refresh the `auth_sig` value by reinstalling the key with a new signature, to prevent replay attacks.

SDE (Store Encrypted)

Instruction mnemonic:

```
sde rs2, offset(rs1)
```

Where:

- `rs1` – base register containing the memory address
- `rs2` – source register whose value will be encrypted and stored
- `offset` – 12-bit signed immediate (added to `rs1`)

Example asm code:

```
sde p3, 8(x10)
```

S-Type RISC-V instruction format:

```
| 31          25 | 24-20 | 19-15 | 14-12 | 11          7 | 6-0 |  
| imm[14:5]  | rs2  | rs1  | 0x1   | imm[4:0] | 0xb |
```

Instruction semantics: When a secret register value is written to memory using `sde`, the following occurs:

1. At decode time, Mojo-V verifies that secret mode is enabled (`mojov_en = 1` in CSR `mojov_cfg`) and that the source register (`rs2`) is a secret register; otherwise, a *Mojo-V security exception* is raised.
2. A new random salt is generated using the processor's true RNG.
3. The plaintext structure (weak or strong, depending on the data contract's `format_sel` field) is populated with the secret register value (`rs2`), salt, authentication signature (`auth_sig`), and any metadata fields defined by the data contract.
4. The plaintext block is encrypted using the symmetric key established through `mojov_keycfg`.
5. The resulting ciphertext is stored to the target memory address.

The encryption of the plaintext data must utilize a **block cipher with full avalanche effect**. As such, any changes in ciphertext will corrupt every bit of the plaintext and vice versa. This process ensures semantic security — identical plaintext values produce unrelated ciphertexts — and authenticity, where any modification to ciphertext causes decryption failure.

LDE (Load Encrypted)

Instruction mnemonic:

```
lde rd, offset(rs1)
```

Where:

- **rs1** – base register containing the memory address
- **rd** – destination register which will receive loaded and decrypted value
- **offset** – 12-bit signed immediate (added to **rs1**)

Example asm code:

```
lde p2, 16(x4)
```

I-Type RISC-V instruction format:

```
| 31          20 | 19-15 | 14-12 | 11-7 | 6-0 |  
| imm[11:0] | rs1 | 0x0  | rd | 0xb |
```

Instruction semantics: When ciphertext is loaded from memory into a secret register using **lde**, the hardware performs the following steps:

1. At decode time, the processor verifies that secret mode is enabled (**mojov_en** = 1 in CSR **mojov_cfg**) and that the destination (**rd**) is a secret register. Violations raise a *Mojo-V security exception*.
2. The ciphertext block is fetched from memory.
3. The block is decrypted using the symmetric key from the current data contract, which is active in the core's load/store encryption unit.
4. The plaintext block's signature (**auth_sig**) is compared against its embedded authentication signature (**auth_sig**). If a mismatch occurs, a security exception is raised and the data is discarded.
5. The plaintext's **val** field is written into the destination secret register (**rd**).

In both load and store paths, the **auth_sig** binds every ciphertext to the current data contract and cryptographic configuration. As a result, ciphertext encrypted under an older contract or a different data-owner key will fail authentication when loaded.

Through the combination of **lde** and **sde** instructions, Mojo-V implements a closed cryptographic datapath where data can only enter and exit the processor through **third-party authenticated encryption**. This mechanism enforces end-to-end confidentiality, integrity, and freshness while eliminating the need to trust any software layer in the computation chain.

5.5 Key Management Interfaces

The Mojo-V key management interface defines how cryptographic keys and **data contracts** are provisioned, maintained, and revoked within the processor. This mechanism forms the foundation for establishing trust between a data owner and the Mojo-V hardware during secret computation.

Mojo-V key management is anchored in hardware and operates entirely within the processor's trusted boundary. No software, including privileged system software or hypervisors, can view or modify plaintext key material. All key provisioning, installation, and verification occur through hardware-controlled interfaces that use the CSRs introduced in Section 4.1.

Data Contract Format. The **data contract** is a 512-bit structure provided by the data owner, encrypted under the processor's **mojov_pubkey**, and installed into the hardware via the **mojov_keycfg** CSR. It defines the symmetric key, format, and authentication parameters used during third-party encrypted computation. The structure aligns on 128-bit boundaries and is defined as:

```
struct { // alignment of 128-bits
    uint8_t sig[16]; // Data contract authentication signature ("Mojo-V ver. #001")
    uint128_t sym_key_128; // 128-bit symmetric key
    uint64_t auth_sig; // 32-bit or 64-bit authentication signature (weak/strong format)
    uint64_t salt; // 64-bit random salt
    uint64_t ciphers; // 64-bit mask indicating asymmetric and symmetric ciphers used
    bool format_sel; // Encryption format selector: 0 = weak, 1 = strong
    uint8_t __pad[]; // Random padding to make total size 512 bits
} data_contract_t ;
```

The **data contract** unifies key material, authentication information, and encryption format selection into a single secure payload. The **auth_sig** value is a signature inserted into all third-party encrypted data, ensuring two security guarantees: *i*) the ciphertext has not been manipulated or forged, and *ii*) no replay of old ciphertexts has occurred since the last data contract installation.

The **ciphers** field of the data contract indicates the asymmetric and symmetric key ciphers that the data owner wishes to use. Only two bits can be set, for one asymmetric and symmetric key cipher. If *i*) more or less bits are set, or *ii*) the hardware doesn't support either of the requested ciphers, then the data contract will not be successfully installed. The format of this field is the same as the format of the Mojo-V **mojov_ciphers** CSR register.

Key Provisioning and Installation. A data owner transmits an encrypted data contract to the Mojo-V processor using secure attestation based on the hardware's **mojov_pubkey**. Once authenticated, the contract is written into the **mojov_keycfg** CSR. This contract contains the

data owner's symmetric key, permitted operations, and the embedded authentication signature (`auth_sig`) used to detect tampering or replay.

Key activation and validation. After installation, the processor validates the data contract. If authentication passes, the `key_valid` bit in the `mojov_cfg` register is set, enabling secret computation. If validation fails, the contract is rejected, and secret mode remains disabled. The hardware enforces that only one valid data contract may be active at any given time.

Key state preservation. During context switches or power transitions, active key state is preserved securely within the `mojov_keystate` CSR. When the key state is saved or restored, it remains encrypted, ensuring that even a malicious operating system cannot observe key material or intermediate values.

Key revocation and freshness. The data owner retains exclusive control of key rotation and revocation. When a data contract is revoked or updated, a new `auth_sig` value is installed to invalidate all previously encrypted data and prevent replay attacks. This ensures that only data encrypted under the current contract can be successfully loaded and decrypted.

Future extensions. A future revision of this specification will expand this section to define the detailed cryptographic protocol for:

- Public key attestation and trust establishment.
- Support for multi-key privacy-oriented applications.
- Hardware-based key derivation for sub-domains and multi-party computation.

These enhancements will further strengthen Mojo-V's capability to operate securely across distributed and multi-tenant environments while maintaining its principle of zero software trust.

6. Microarchitectural Considerations

Mojo-V's architecture enforces secrecy and data-oblivious behavior not only through its instruction semantics but also through carefully constrained microarchitectural design. All secret computation rules are applied at the earliest possible stage of execution, ensuring that no transient behavior, speculation, or optimization can leak secret information.

Decode-Time Enforcement. Security enforcement begins at instruction decode. Each instruction is analyzed to determine whether its operands or destinations involve secret registers. The decode stage ensures that secret and non-secret operands are never mixed in unsafe ways, and any violation — such as a secret register used in a branch predicate, load/store address, or CSR modification — immediately triggers a *Mojo-V security exception*. If any of the remaining rules apply, the decode stage also **propagates a single bit that indicates whether the instruction is secret or public**, based on the destination register. Because operand secrecy is known statically at decode, this enforcement is deterministic and leaves no ambiguity for later pipeline stages.

Constant-Latency Functional Units. Mojo-V prevents timing-based information leaks by ensuring that operations involving secret operands execute with timing that is **independent of secret inputs**. Functional units need not be truly constant-time; they must simply ensure that execution latency does not vary with secret data. This constraint applies only to operations explicitly marked as secret (based on the decode-stage secrecy bit). This approach balances performance and security by isolating timing protections to instructions that actually process secret data.

Speculation Safety. Unconstrained speculative execution is permitted in Mojo-V, but all speculation involving secret data is constrained in one easy-to-identify condition. Wherever dependencies are speculatively inserted — primarily within the load/store queue — secret-to-non-secret speculative forwarding must be prohibited, as it could permit Spectre-style leakage of secret data into observable non-secret speculative register state. Secret-to-secret, non-secret-to-secret, and secret-to-secret speculative forwarding are allowed, since they do not expose secrets across trust boundaries. These checks are simple to implement since the decode stage marks every operation unambiguously as secret or non-secret. This model preserves the performance benefits of speculation while preventing secret data from contaminating untrusted program state.

Exception Masking. When instructions involving secret operands trigger exceptions (for example, divide-by-zero or invalid operation), Mojo-V masks those exceptions so that no information about the secret operands is revealed. Masked exceptions neither raise traps to software nor alter observable architectural state. Instead, they are dropped or they propagate an exception metadata bit internally, maintaining silent and data-oblivious behavior. This ensures that secret-related faults cannot act as a side channel.

Compatibility and Integration Requirements. To maintain portability and predictable semantics, Mojo-V remains compliant with the standard RV64 RISC-V instruction set. In addition, Mojo-V implementations must include the **Zicond** conditional-move extension to allow data-oblivious predication. To support compatibility across developments, the **mojov_ver** field of the **mojov_cfg** CSR register includes an 8-bit version number that reflects the precise Mojo-V semantics implemented by the processor. This enables future revisions of the architecture to evolve while preserving backward compatibility and verifiable behavior.

7. System Software Considerations

This section details a number of system software considerations that must be addressed to fully utilize Mojo-V capabilities.

Compiler Support: The Mojo-V project is developing a Mojo-V capable version of the LLVM compiler. The compiler explicitly supports Mojo-V's **secret computation model** by allowing programmers to declare variables and computations as secret. By convention, languages should express this using a **secret** attribute or type qualifier. When a secret variable is

computed upon, the compiler must ensure that all intermediate values remain confined to **secret registers** (p0-p3, pf0-pf3). Any attempt to use a secret value as a control predicate, code pointer, or data pointer must trigger a compile-time error. The compiler should also expose a built-in conditional move intrinsic, such as `mojov_cmov()`, to allow side-channel-free predicated decision-making.

As an example of Mojo-V enabled code, please consider the code in Figure 2, which implements a data-oblivious version of bubble-sort for Mojo-V secret computation. The `data` variable is an array of secret integers. The bubble-sort algorithm iterates over the `size` integers in the array, performing a data-oblivious swap operation with the built-in compiler primitive `mojov_cmov`. `mojov_cmov` performs a data-oblivious conditional move: based on the first secret predicate operand, the primitive outputs the second (if true) or third (if false) operand. As this bubble-sort algorithm runs, the software and developers cannot see or infer the values of the `data` array, regardless of any coding changes or hacking of the bubble-sort software.

```
void bubblesort(secret int *data, unsigned size) {
    for (unsigned i = 0; i < size - 1; i++) {
        for (unsigned j = 0; j < size - 1; j++) {
            // swap needed?
            secret bool swap = (data[j] > data[j + 1]);

            // perform the swap
            secret int tmp = data[j];
            data[j] = mojov_cmov(swap, data[j + 1], data[j]);
            data[j+1] = mojov_cmov(swap, tmp, data[j + 1]);

            // count the number of swaps executed
            swaps = mojov_cmov(swap, swaps + 1, swaps);
        }
    }
}
```

Figure 2: Mojo-V version of the Bubble-sort Algorithm.

There is **no trust placed in the compiler**—if it accidentally emits an instruction that would cause a side channel using a secret value, the Mojo-V hardware will detect this at decode time and raise a Mojo-V *security exception*.

Operating System Support: The operating system must preserve secret state across context switches. When secret mode is active (`mojov_en = 1`), the OS should save and restore secret registers using the **LDE** and **SDE** instructions. However, this does **not** imply trust in the OS—these operations occur entirely on encrypted data. Without OS support, an attempt to context-switch while in secret mode will result in a hardware Mojo-V *security exception* when the OS tries to save the secret registers (e.g., using `sd`), as Mojo-V forbids non-encrypted access to secret register values.

System Libraries: System libraries must be aware of Mojo-V's use of secret registers. Two key implications follow:

(a) **Compatibility builds.** Libraries that run on standard RISC-V systems may use all registers. When building a program with the `-fmojov` compiler flag, the toolchain should link against library versions that **do not** use the secret registers. This separation allows Mojo-V programs to reserve the secret register file without modification to library code. This register partitioning simplifies hardware, avoiding the need for data tagging, and allows all enforcement to occur cleanly at the decode stage.

(b) **Data-oblivious library variants.** Certain libraries—particularly `string`, `math`, and `sorting`—should include Mojo-V-enabled variants that use secret registers and implement **data-oblivious algorithms**. These functions can then safely operate on encrypted data while maintaining hardware-enforced privacy guarantees.

Client and Server Side Mojo-V Libraries: Both **data owners** (clients) and **service providers** (servers) benefit from Mojo-V software support libraries.

- **Client-side libraries** should include utilities for software-based key generation, encryption, and decryption, allowing data owners to prepare encrypted inputs and recover encrypted results.
- **Server-side libraries** should provide data-oblivious math and sorting routines, bindings for higher-level languages such as Python or JavaScript, and secure network marshalling functions to handle encrypted data exchanges.

Mojo-V Debug Support: Debugging in Mojo-V must preserve the same security principles as execution. The hardware must never expose decrypted data through debug interfaces. Instead, developers may use **software-based debugging modes** by inserting a **decrypted data contract** directly into the server environment. This enables debugging tools (such as GDB) to inspect data under development keys using software-based decryption algorithms. Production workloads using third-party keys cannot be debugged in plaintext—this is a feature, not a limitation, of Mojo-V's design. This guarantee ensures that even during development, production secrets remain inaccessible to software, administrators, or developers.

8. Authorship

Primary author: Todd Austin, University of Michigan / Agita Labs, austin@umich.edu

Additional contributors:

- Lauren Biernacki, Lafayette College
- Shibo Chen, Tenstorrent
- Meron Demissie, University of Michigan
- Yonathan Fisseha, University of Michigan
- Tarunesh Verma, University of Michigan

9. Change Log

Version 0.90 – October 24, 2025, Initial release.

Appendix A. Comparison to FHE

Both **Fully Homomorphic Encryption (FHE)** and **Mojo-V** enable computation over encrypted data, but they differ fundamentally in their approach, performance, and security properties.

1. **Mathematical vs. Architectural Solutions**

FHE is a **mathematical** solution to secret computation, built entirely upon cryptographic algebra and lattice-based security proofs. Mojo-V, by contrast, is an **architectural** solution—achieving secret computation by embedding secrecy enforcement directly into the processor hardware. In FHE, privacy derives from mathematical intractability; in Mojo-V, it derives from physical isolation and microarchitectural enforcement.

2. **Programming Experience and Performance**

To programmers, both systems appear similar—computation proceeds directly on encrypted data, and a memory dump or software compromise reveals only ciphertext. However, Mojo-V executes using standard processor hardware and runs roughly **10 million times faster** than a comparable CPU-based hybrid integer–Boolean FHE implementation. This makes Mojo-V the first architecture to deliver *interactive, human-scale secret computation* in real time.

3. **Attack Model and Threat Surface**

From an attacker’s perspective, FHE provides stronger theoretical protection, as it remains secure even against physical inspection. Mojo-V, on the other hand, is susceptible to **analog measurement attacks**—such as observing CPU power, thermal, or voltage fluctuations—that could potentially leak secret information. Because of this distinction, Mojo-V is characterized as a **consumer-grade privacy technology**, while FHE represents a **military-grade privacy technology** with provable cryptographic guarantees. A key aspect of our follow-on work is to harden the plaintext portion of the

Mojo-V hardware implementation against differential and direct analog measurement attacks, which will serve to narrow the security gap between Mojo-V and FHE.

4. **Data Representation and Programming Style**

FHE systems generally offer two computational paradigms: **Integer FHE**, which is faster but difficult to express arbitrary programs in, and **Boolean FHE**, which is easier to program and fully Turing complete. Mojo-V's programming model is conceptually closer to Boolean FHE—supporting arbitrary control structures, loops, and data types—but operates at **near-native hardware speed**, vastly outperforming even optimized integer FHE platforms. Unlike FHE, Mojo-V supports standard floating-point modes.

5. **Roadmap and Future Integration**

While Mojo-V does not rely on mathematical homomorphism, future versions will integrate **concise proofs** and **safe disclosures**—mechanisms similar to zero-knowledge proofs (ZKPs)—to allow verifiable results without revealing secret data. This will bridge Mojo-V's architectural secrecy with cryptographic proof-of-correctness, providing hybrid trust assurance between physical enforcement and mathematical verifiability.

In summary, **FHE secures data through pure mathematics**, while **Mojo-V secures data through hardware semantics**. Both protect against software-level threats, but Mojo-V delivers practical performance and programmability for consumer and enterprise workloads, complementing FHE's formal cryptographic strength.

Appendix B. Comparison to CHERI

The **CHERI (Capability Hardware Enhanced RISC Instructions)** architecture and **Mojo-V** share the goal of improving system trustworthiness through hardware support, but they approach the problem from fundamentally different perspectives.

CHERI extends a conventional RISC instruction set, compiler, and operating system to implement **fine-grained, capability-based memory protection**. Its primary goal is to mitigate **memory-safety vulnerabilities** in unsafe languages such as C by introducing unforgeable capabilities—pointers with explicit bounds and permissions—that prevent spatial and temporal memory errors. By combining these capabilities with operating-system-managed object compartments, CHERI provides **software compartmentalization** that can contain and limit damage from buffer overflows, pointer corruption, and similar exploit classes.

In contrast, Mojo-V ignores every vulnerability that CHERI tries to stop, and it instead **takes a data-centric approach**. Mojo-V does not attempt to eliminate software vulnerabilities; instead, it renders them harmless by protecting the *data* that attackers seek. Mojo-V uses **cryptographic isolation** and **hardware-enforced secret registers** to ensure that even if an attacker fully exploits a vulnerability, all retrieved information remains encrypted. While CHERI treats the *symptoms*—preventing corruption and unauthorized access through software correctness and pointer control—Mojo-V treats the *disease* by eliminating any value in the data exposed to attackers.

Conceptual Comparison:

Dimension	CHERI	Mojo-V
Security Objective	Prevents software vulnerabilities and enforces memory safety	Protects secret data from disclosure even in the presence of vulnerabilities
Primary Mechanism	Hardware-enforced capabilities for pointers and memory bounds	Hardware-enforced secrecy domain with encrypted storage and secret registers
Software Trust Model	Requires trusted compiler and OS to enforce capability semantics	Zero software trust; hardware enforces secrecy and integrity autonomously
Granularity of Protection	Per-pointer and per-object memory bounds	Per-register and per-memory-block cryptographic isolation
Failure Model	Compromise prevented by safety rules	Compromise tolerated but yields only ciphertext
Performance Impact	Low to moderate overhead for capability manipulation	Negligible overhead—executed at native CPU speed
Hardware Extension	Adds capability registers and tagged memory	Adds secret registers, encryption CSRs, and data-contract mechanism

Broader Security Impact: Mojo-V also mitigates a range of *microarchitectural* and *physical-layer* exploits that CHERI does not address. Because CHERI preserves the classic memory-based computation model, it remains susceptible to many side-channel and speculative-execution attacks that exploit timing, cache, or speculation leakage. Mojo-V, by contrast, **sequesters secret computation** and **encrypts all externalized state**, preventing data exposure from both digital and analog vectors.

Mojo-V neutralizes attacks beyond CHERI's scope, including:

- Rowhammer
- Spectre V1/V2 and related speculative attacks
- Portsmash and Foreshadow
- Microarchitectural Data Sampling (MDS, RIDL, Zombieload, Fallout)
- Cold-boot and DMA extraction attacks
- PRIME+PROBE and other cache timing attacks
- Stack canary escapes
- GoFetch and prefetcher leakage
- Ciphertext analysis and CIPHERLEAK
- Kocher and Bernstein timing side-channels
- NBTI-induced leakage and aging-based disclosure

- Malicious developer or insider exfiltration
- And likely many *future, not-yet-invented* attack classes

Summary: CHERI provides a **software-centric solution** to software vulnerability mitigation—an evolutionary extension of the RISC model that enforces safer pointers and controlled compartmentalization through a trusted compiler and operating system. Mojo-V, in contrast, provides a **data-centric, cryptographically rooted solution** that assumes software may be compromised and instead ensures that no secret information can ever be exposed. In essence, **CHERI hardens code**, while **Mojo-V hardens data**, offering a simpler, faster, and more future-proof approach to computational privacy.